

PAC-Bayesian Generalization Bounds for Weight-Decomposed Low-Rank Adaptation

Stephen Zhang

Varun Prakash

Soham Vasanawala

June 3, 2026

Abstract

Parameter-efficient fine-tuning (PEFT) methods such as DoRA and MAP decompose weight updates into magnitude and directional components, yet their generalization behavior is poorly understood. We give a PAC-Bayesian analysis of weight-decomposed adaptation using von Mises–Fisher (vMF) priors on the hypersphere, yielding closed-form KL complexity terms in the angular deviations between pretrained and fine-tuned directions. We prove three results and verify each on RoBERTa-base fine-tuned across six GLUE tasks: a *non-asymptotic directional KL bracket* (Theorem 3.4), valid beyond the high-concentration regime; an *angular complexity comparison* (Theorem 3.10) showing that column-wise methods (DoRA) *sum* directional complexity over local angles while global methods (MAP) *average* it into a single rotation; and a *sparse-regime advantage* (Theorem 3.12) characterizing when each method is preferable. Our verification targets the geometric complexity quantities entering the bounds (angular KL, the comparison identity, effective sparsity), not realized test risk. Measuring per-column angles over 2592 weight matrices, the comparison identity holds to machine precision (and with $R^2=0.995$ in its measured form), the KL bracket holds with zero violations and a $77\text{--}307\times$ asymptotic-to-non-asymptotic tightening, and the effective sparsity s/d cleanly separates a sparse, DoRA-favored regime (RTE and MRPC, $s/d\approx 0.10$) from dense, MAP-favored regimes (SST-2/QNLI/MNLI/QQP, $s/d\geq 0.95$), as the theory predicts.

1 Introduction

Large models are adapted to downstream tasks by parameter-efficient fine-tuning (PEFT), which modifies a small subset of parameters. DoRA [1] updates each *column* of a weight matrix independently (column-wise rotations); MAP [2] rotates the entire flattened matrix as one vector. There is no theory explaining when one generalizes better than the other. We close this gap with a PAC-Bayesian analysis: modeling directional updates with vMF distributions on the hypersphere yields bounds depending explicitly on angular deviations, giving a notion of *directional complexity*. We prove a non-asymptotic KL bound (Theorem 3.4), an angular comparison theorem (Theorem 3.10), and a sparse-regime advantage result (Theorem 3.12), and verify all three empirically on RoBERTa-base.

2 Background

PEFT. LoRA [3] writes $\Delta W = BA$, $B \in \mathbb{R}^{d \times r}$, $A \in \mathbb{R}^{r \times d}$, $r \ll d$. DoRA decomposes each column $w_j = m_j u_j$, $m_j \in \mathbb{R}_+$, $u_j \in \mathbb{S}^{d-1}$, updating the direction via LoRA and the magnitude directly. MAP flattens $W \in \mathbb{R}^{d \times d}$ to $w \in \mathbb{R}^{d^2}$ and writes $w' = \alpha \hat{w} + \beta \widehat{\Delta w}$ with unit $\hat{w}, \widehat{\Delta w}$ and learnable α, β .

PAC-Bayes. Bounds [4] control the risk of a randomized predictor through $\text{KL}(Q\|P)$ between posterior Q and data-free prior P .

vMF. vMF(μ, κ) on $\mathbb{S}^{d-1}$ has density $p(x \mid \mu, \kappa) = C_d(\kappa) \exp(\kappa \mu^\top x)$ with mean resultant length $A_d(\kappa) = I_{d/2}(\kappa)/I_{d/2-1}(\kappa) \in [0, 1]$, increasing in κ .

3 PAC-Bayesian Analysis of Weight-Decomposed Adaptation

Let P be a data-free prior and Q a posterior learned from $S \sim \mathcal{D}^n$, loss $\ell \in [0, 1]$, with $R(Q) = \mathbb{E}_{h \sim Q} \mathbb{E}_z[\ell]$ and $\hat{R}(Q) = \mathbb{E}_{h \sim Q} \frac{1}{n} \sum_i \ell(h, z_i)$.

Theorem 3.1 (McAllester [4]). *For any $\delta \in (0, 1)$, w.p. $\geq 1 - \delta$, for all Q : $R(Q) \leq \hat{R}(Q) + \sqrt{(\text{KL}(Q\|P) + \ln \frac{2\sqrt{n}}{\delta})/2n}$.*

Bounding $R(Q)$ thus reduces to bounding $\text{KL}(Q\|P)$ under the structured parameterizations of DoRA and MAP.

Lemma 3.2 (vMF KL, equal concentration). *For $Q_1 = \text{vMF}(\mu_1, \kappa)$, $Q_0 = \text{vMF}(\mu_0, \kappa)$ on $\mathbb{S}^{d-1}$: $\text{KL}(Q_1\|Q_0) = \kappa A_d(\kappa)(1 - \mu_1^\top \mu_0)$.*

Proof. $\text{KL} = \mathbb{E}_{Q_1}[\ln p(x | \mu_1, \kappa) - \ln p(x | \mu_0, \kappa)] = \kappa(\mu_1 - \mu_0)^\top \mathbb{E}_{Q_1}[x]$, since the normalizers cancel at equal κ . The vMF mean is $\mathbb{E}_{Q_1}[x] = A_d(\kappa)\mu_1$, so $\text{KL} = \kappa A_d(\kappa)(1 - \mu_1^\top \mu_0)$. $\square$

Lemma 3.3 (Uniform bounds on $A_d(\kappa)$). *For all $\kappa > 0$, $d \geq 2$: $\kappa/(d + \kappa) \leq A_d(\kappa) \leq 1$.*

Theorem 3.4 (Non-Asymptotic Directional KL Bound). *Under the column-wise decomposition of DoRA with per-column vMF priors of concentration κ and angles θ_j ,*

$$\frac{\kappa^2}{d + \kappa} \sum_{j=1}^d (1 - \cos \theta_j) \leq C_{\text{dir}}^{\text{DoRA}} \leq \kappa \sum_{j=1}^d (1 - \cos \theta_j),$$

and under the global decomposition of MAP with angle Θ_{global} ,

$$\frac{\kappa^2}{k + \kappa} (1 - \cos \Theta_{\text{global}}) \leq C_{\text{dir}}^{\text{MAP}} \leq \kappa (1 - \cos \Theta_{\text{global}}).$$

Proof. By Lemma 3.2, $C_{\text{dir}}^{\text{DoRA}} = \sum_j \text{KL}(Q_{\text{dir},j} \| P_{\text{dir},j}) = \kappa A_d(\kappa) \sum_j (1 - \cos \theta_j)$. Substituting the two bounds on $A_d(\kappa)$ from Lemma 3.3 gives $\frac{\kappa^2}{d + \kappa} \sum_j (1 - \cos \theta_j) \leq \kappa A_d(\kappa) \sum_j (1 - \cos \theta_j) \leq \kappa \sum_j (1 - \cos \theta_j)$. The MAP case is identical with $d \rightarrow k$ (the flattened dimension) and the sum replaced by the single global term. $\square$

Remark 3.5. *The bracket holds for all $\kappa > 0$, removing any need for $\kappa \rightarrow \infty$. Since $A_d(\kappa) \approx \kappa/d$ for $\kappa \ll d$, the exact KL sits near the lower bound, and the asymptotic approximation ($A_d \equiv 1$, the upper bound) overestimates the directional KL by a factor $1/A_d(\kappa)$ —quantified in section 4.*

3.1 Local vs. Global: The Comparison Theorem

For DoRA, the directional complexity is $C_{\text{dir}}^{\text{DoRA}} = \kappa A_d(\kappa) \sum_j (1 - \cos \theta_j)$ with $\theta_j = \arccos(\hat{\mu}_j^\top \mu_{0,j})$ (the per-column angle); for MAP it is $\kappa A_k(\kappa)(1 - \cos \Theta_{\text{global}})$ with Θ_{global} the angle of the single global rotation. The next theorem relates the two.

Theorem 3.10 (Angular Complexity Comparison). *Let $u_j, \hat{u}_j \in \mathbb{S}^{d-1}$ be the pretrained and fine-tuned unit directions for column j , with $\theta_j \in [0, \pi/2]$. Define the concatenated unit vectors $U = \frac{1}{\sqrt{d}}(u_1, \dots, u_d)$, $\hat{U} = \frac{1}{\sqrt{d}}(\hat{u}_1, \dots, \hat{u}_d)$. Then $\Theta_{\text{global}} = \arccos(\hat{U}^\top U)$ satisfies*

$$\Theta_{\text{global}}^2 \leq \frac{1}{d} \sum_{j=1}^d \theta_j^2 \leq \sum_{j=1}^d \theta_j^2, \quad (1)$$

and the following identity holds exactly:

$$1 - \cos \Theta_{\text{global}} = \frac{1}{d} \sum_{j=1}^d (1 - \cos \theta_j). \quad (2)$$

Proof. Step 1 (exact identity). From the definitions, $\cos \Theta_{\text{global}} = \hat{U}^\top U = \frac{1}{d} \sum_j \hat{u}_j^\top u_j = \frac{1}{d} \sum_j \cos \theta_j$, which rearranges to (2). **Step 2 (convexity of $\arccos$).** Set $c_j = \cos \theta_j$ and $h(c) = (\arccos c)^2$. Writing $\theta = \arccos c \in [0, \pi/2]$, $h''(c) = \frac{2}{\sin^2 \theta} (1 - \theta \cot \theta) \geq 0$ since $\theta \leq \tan \theta$ on $[0, \pi/2]$, so h is convex on $[0, 1]$. By (2) ($\cos \Theta_{\text{global}} = \frac{1}{d} \sum_j c_j$) and Jensen's inequality applied to h , $\Theta_{\text{global}}^2 = h(\frac{1}{d} \sum_j c_j) \leq \frac{1}{d} \sum_j h(c_j) = \frac{1}{d} \sum_j \theta_j^2 \leq \sum_j \theta_j^2$, the last step as $d \geq 1$. (Note $\cos$ is *concave* on $[0, \pi/2]$, so the bound follows from convexity of $c \mapsto (\arccos c)^2$, not from Jensen on $\cos$ directly.) $\square$

Remark 3.11. *The global angle is an average over local angles (in the sense of averaging cosines), whereas $C_{\text{dir}}^{\text{DoRA}}$ is a sum: DoRA sums what MAP averages. MAP thus always has a smaller-or-equal global angular deviation, at the cost of per-column resolution.*

3.2 When Does DoRA Achieve a Tighter Bound?

Theorem 3.12 (Localized-Update Complexity Bounds). *Suppose the update is sparse: for some $\mathcal{S} \subseteq \{1, \dots, d\}$ with $|\mathcal{S}| = s \ll d$, $\theta_j = 0$ for $j \notin \mathcal{S}$ and $\theta_j \leq \varepsilon$ for $j \in \mathcal{S}$. Then*

$$C_{\text{dir}}^{\text{DoRA}} \leq \kappa A_d(\kappa) s (1 - \cos \varepsilon), \quad C_{\text{dir}}^{\text{MAP}} \leq \kappa A_k(\kappa) \frac{s}{d} (1 - \cos \varepsilon).$$

Proof. For DoRA, $\theta_j = 0$ off $\mathcal{S}$ contributes zero, and $1 - \cos \theta_j \leq 1 - \cos \varepsilon$ for $j \in \mathcal{S}$, giving $C_{\text{dir}}^{\text{DoRA}} = \kappa A_d(\kappa) \sum_{j \in \mathcal{S}} (1 - \cos \theta_j) \leq \kappa A_d(\kappa) s (1 - \cos \varepsilon)$. For MAP, the exact identity (2) with $\theta_j = 0$ off $\mathcal{S}$ gives $1 - \cos \Theta_{\text{global}} = \frac{1}{d} \sum_{j \in \mathcal{S}} (1 - \cos \theta_j) \leq \frac{s}{d} (1 - \cos \varepsilon)$; substituting into the upper bound of Theorem 3.4 yields the claim. $\square$

Remark 3.13 (Sparse vs. dense; why raw complexity cannot decide). *The identity (2) gives $\sum_j (1 - \cos \theta_j) = d(1 - \cos \Theta_{\text{global}})$ exactly, hence $C_{\text{dir}}^{\text{DoRA}} = d C_{\text{dir}}^{\text{MAP}}$ for every update: MAP’s raw directional KL is always a factor $\sim d$ smaller and raw complexity never crosses over between regimes. Consequently the raw bounds of Theorem 3.12 favor MAP in every regime; the DoRA advantage is a statement about dimension-calibrated discriminability (below), not about the raw upper bounds. The meaningful comparison is on dimension-calibrated complexity—scaling $\kappa_{\text{eff}} = \kappa d$ for DoRA (on $\mathbb{S}^{d-1}$) and $\kappa_{\text{eff}} = \kappa k$ for MAP (on $\mathbb{S}^{k-1}$), which cancels this $1/d$ since $A_n(\kappa_{\text{eff}}) \approx \kappa$ for either sphere. What remains is DoRA’s ability to allocate complexity non-uniformly across columns (its discriminability). That ability is valuable when updates are sparse (few nonzero θ_j ; DoRA gives an $O(s)$ description resolving which columns moved) and wasted when updates are dense ($\theta_j \approx \Theta_{\text{global}}$ uniformly; a single rotation suffices). Hence DoRA is favored in the sparse regime, MAP in the dense one.*

4 Empirical Verification

The theorems are proven; we now verify that the geometric quantities they rest on behave as claimed on real fine-tuned weights. We do not, in this draft, evaluate the realized McAllester bound (Theorem 3.1) end-to-end; a directional-only estimate of its magnitude is reported at the end of section 4.2 (the magnitude-component KL is left to future work). We fine-tune RoBERTa-base (125M) [5] under LoRA, DoRA, and MAP (rank $r=8$, matched parameter count, $\kappa=10$) on six GLUE tasks (SST-2, QNLI, MNLI, RTE, QQP, MRPC; three seeds, full training sets), and for each of the 72 target weight matrices per run measure the per-column angles $\theta_j = \arccos(w_{0,j}^\top \hat{w}_j / (\|w_{0,j}\| \|\hat{w}_j\|))$ and the global angle Θ_{global} between pretrained W_0 and merged fine-tuned $\hat{W} = W_0 + \Delta W$. Pooling the four primary tasks gives $72 \times 36 = 2592$ measured layers; QQP and MRPC add 1296 more. We evaluate $\hat{C}_{\text{dir}}^{\text{DoRA}} = \kappa A_d(\kappa) \sum_j (1 - \cos \theta_j)$ exactly as in the theory, with $A_d(\kappa) = I_{d/2}(\kappa) / I_{d/2-1}(\kappa)$ computed by an Olver large-order expansion to avoid Bessel underflow at $d = 768, 3072$.

4.1 Verifying Theorem 3.10 (sum-vs-average identity)

In its exact unit-column form, the identity (2) holds to machine precision on the measured angles (max residual 2.2×10^{-16}). In the form the pipeline instruments—raw, non-unit column flattening for Θ_{global} —regressing $d(1 - \cos \Theta_{\text{global}})$ on $\sum_j (1 - \cos \theta_j)$ across all 2592 layers gives slope 0.889 with $R^2 = 0.995$ (Pearson $r = 0.997$); fig. 1(left) (the high R^2 is expected, since both axes are the same quantity up to the flattening offset—the load-bearing check is the 2.2×10^{-16} unit-column residual above, not the regression fit). The near-perfect linearity confirms DoRA’s local sum is the dimension-scaled counterpart of MAP’s global angle; the 11% slope shortfall is consistent with the raw-vs-unit flattening offset (the raw global cosine is a magnitude-weighted average of $\cos \theta_j$), not a violation. The held-out QQP and MRPC runs (1296 further layers) reproduce the fit (slope 0.866, $R^2 = 0.995$, Pearson $r = 0.998$). Table 1 shows the resulting complexity terms, with $C^{\text{DoRA}} / C^{\text{MAP}} \approx 10^3$ on every task—the sum-vs-average gap of Theorem 3.10 made quantitative.

4.2 Verifying Theorem 3.4 (non-asymptotic KL bracket)

The bracket $\frac{\kappa^2}{d+\kappa} S \leq \kappa A_d(\kappa) S \leq \kappa S$ ($S = \sum_j (1 - \cos \theta_j)$) holds with **0 violations across all 2592 layers** (fig. 1, right; and 0/1296 on the held-out QQP/MRPC layers). Because the bracket follows algebraically from Lemma 3.3 and $S \geq 0$, it cannot be violated by construction; the zero-violation count instead confirms that the Olver large-order expansion for $A_d(\kappa)$ stays numerically stable at $d = 768, 3072$ (no Bessel underflow), and the substantive content is the $1/A_d(\kappa)$ tightening below. The exact value sits at the lower end of the bracket (0.01% of the way to the upper bound) because $A_d(\kappa) \ll 1$ at these dimensions: $A_{768}(10) = 1.30 \times 10^{-2}$, $A_{3072}(10) = 3.26 \times 10^{-3}$. Equivalently (Remark 3.5), the asymptotic approximation overestimates the directional KL by $1/A_d(\kappa) \approx \mathbf{77} \times$ ($d=768$) and $\mathbf{307} \times$ ($d=3072$): MNLI/DoRA’s directional KL falls from 3903 (asymptotic) to 46.3 (non-asymptotic).

Realized bound (directional-only). Plugging the measured directional KL into McAllester’s bound (Theorem 3.1) gives a non-vacuous complexity term $\sqrt{(\text{KL} + \ln \frac{2\sqrt{n}}{\delta})/2n}$ ($\delta=0.05$): for MNLI/DoRA (KL=46.3, $n=392,702$) it is 0.0085, for QQP/DoRA (KL=26.9, $n=363,846$) 0.0071, and for MRPC/DoRA (KL=0.21, $n=3,668$) 0.033—in every case well below 1. Two caveats. (i) This uses the *directional* KL only; the magnitude-component KL (deferred to future work) only increases KL, so this is a *lower bound* on the true complexity term. (ii) On raw directional KL $C^{\text{DoRA}} \approx d C^{\text{MAP}}$ (Remark 3.13), so this must not be read as a DoRA-vs-MAP bound comparison—the DoRA advantage is in dimension-calibrated discriminability (section 3.2), not the raw realized bound.

4.3 Verifying Theorem 3.12 (sparse $\Rightarrow$ DoRA, dense $\Rightarrow$ MAP)

What not to measure. A tempting test is a crossover in the raw complexity ratio $C_{\text{dir}}^{\text{DoRA}} / C_{\text{dir}}^{\text{MAP}}$. But Remark 3.13 shows the identity (2) forces this ratio to equal d *identically*, in every regime; our data confirm it—the measured ratio is




Figure 1: **Verification on 2592 measured layers** (four primary tasks; QQP/MRPC add 1296 more). (Left) Theorem 3.10: local sum vs. dimension-scaled global term in the *raw* (non-unit) flattening, near $y = x$ ($R^2=0.995$, slope=0.89); the exact unit-column identity itself holds to 2.2×10^{-16} . (Right) Theorem 3.4: the exact $C_{\text{dir}} = \kappa A_d(\kappa)S$ (red) lies inside the analytic bracket (shaded) for a representative MNLI/DoRA run; 0/2592 violations overall, hugging the lower bound since $A_d(\kappa) \ll 1$ (the bracket holds by construction—the panel shows numerical stability and the size of the asymptotic-to-exact gap).

Table 1: Measured directional complexity (summed over 72 layers, mean over 3 seeds, $\kappa=10$) and effective sparsity s/d (LoRA/DoRA). The $\sim 10^3$ ratio $C^{\text{DoRA}}/C^{\text{MAP}}$ is the sum-vs-average gap (Theorem 3.10); s/d separates the sparse (RTE, MRPC) and dense (others) regimes of Theorem 3.12. QQP and MRPC (below the rule) are held-out tasks.

Task	Method	C^{DoRA} (asym)	C^{DoRA} (non-asym)	C^{MAP}	$C^{\text{DoRA}}/C^{\text{MAP}}$	s/d
RTE	LoRA	17.1	0.21	0.016	1066	0.098
	DoRA	17.0	0.20	0.017	1008	0.103
SST-2	LoRA	498.1	5.76	0.476	1047	0.935
	DoRA	520.0	6.03	0.499	1043	0.964
QNLI	DoRA	969.8	11.56	0.864	1122	0.986
MNLI	LoRA	3837.9	45.43	3.461	1109	0.998
	DoRA	3902.9	46.30	3.528	1106	1.000
QQP	LoRA	2255.7	26.83	1.838	1227	0.990
	DoRA	2264.3	26.93	1.846	1227	0.998
MRPC	LoRA	16.8	0.20	0.017	1007	0.105
	DoRA	17.9	0.21	0.018	1009	0.110

1009, 1043, 1122, 1106 for RTE/SST-2/QNLI/MNLI (and 1227, 1009 for QQP/MRPC), i.e. $\approx d$ regardless of sparsity. Raw complexity therefore cannot distinguish the regimes, consistent with Theorem 3.12/Remark 3.13. The regime-discriminating quantity is the *dimension-calibrated discriminability*: holding κ_{eff} fixed, how *non-uniformly* DoRA’s complexity is distributed across columns. We quantify this by the concentration $\mathcal{D} = (\text{fraction of } \sum_j(1 - \cos \theta_j) \text{ in active columns } \theta_j > \tau) / (s/d)$, the ratio of complexity mass captured to the column budget spent; $\mathcal{D}=1$ is a perfectly uniform (dense) update and $\mathcal{D} \gg 1$ a concentrated (sparse) one. By construction $\mathcal{D}$ is large precisely when directional complexity concentrates in few columns, i.e. when the update is sparse; it therefore measures the update *geometry* (the precondition of Theorem 3.12) rather than an independently realized DoRA bound. We report it as evidence that the tasks fall into the predicted regimes, not as a direct measurement of a tighter DoRA bound.

Result. The effective sparsity $s/d = |\{j : \theta_j > \tau\}|/d$ ($\tau=0.01$ rad) and the concentration $\mathcal{D}$ separate the six tasks cleanly into the two predicted regimes (table 2):

- **Sparse $\Rightarrow$ DoRA-favored.** RTE and MRPC ($s/d \approx 0.10$): only $\sim 10\%$ of columns move, yet they carry $\sim 30\%$ of the directional complexity, a concentration $\mathcal{D} \approx 3.0$. DoRA’s per-column $O(s)$ description resolves which columns changed and allocates complexity there; MAP collapses this localized update into one angle and cannot distinguish it from a diffuse one. DoRA is the more discriminating measure, as Theorem 3.12 predicts.
- **Dense $\Rightarrow$ MAP-favored.** SST-2, QNLI, MNLI, QQP ($s/d = 0.95\text{--}1.00$): nearly all columns rotate and complexity

Table 2: Regime separation for Theorem 3.12 (LoRA/DoRA, mean over layers and seeds). Raw $C^{\text{DoRA}}/C^{\text{MAP}} \approx d$ in all regimes (it cannot cross over); the dimension-calibrated discriminability $\mathcal{D}$ does separate them.

Task	s/d	mass in active cols	$\mathcal{D}$	regime / favored method
RTE	0.10	0.30	3.00	sparse $\Rightarrow$ DoRA
MRPC	0.11	0.33	3.10	sparse $\Rightarrow$ DoRA
SST-2	0.95	0.99	1.05	dense $\Rightarrow$ MAP
QNLI	0.98	1.00	1.02	dense $\Rightarrow$ MAP
QQP	0.99	1.00	1.01	dense $\Rightarrow$ MAP
MNLI	1.00	1.00	1.00	dense $\Rightarrow$ MAP

is spread uniformly, $\mathcal{D} \approx 1.0$. DoRA’s column-wise flexibility buys nothing the single MAP rotation does not already capture, so MAP attains the same effect with one calibrated degree of freedom and is preferred.

The discriminability $\mathcal{D}$ thus decreases from ≈ 3.0 (the sparse tasks RTE, MRPC) to ≈ 1.0 (the dense tasks SST-2/QNLI/MNLI/QQP) as s/d rises from 0.10 to 1.00: the sparse $\rightarrow$ dense, DoRA $\rightarrow$ MAP transition of Theorem 3.12 is realized across tasks on the calibration-consistent quantity, not on raw complexity.

5 Conclusion

We proved a non-asymptotic directional KL bound (Theorem 3.4), an angular complexity comparison (Theorem 3.10), and localized-update complexity bounds (Theorem 3.12) for weight-decomposed PEFT, and verified the geometric quantities on 2592 measured RoBERTa-base layers (plus 1296 held-out): the sum-vs-average identity holds to machine precision, the KL bracket holds with zero violations and a 77–307 $\times$ non-asymptotic tightening, and effective sparsity separates sparse (RTE, MRPC) from dense (SST-2/QNLI/MNLI/QQP) regimes—the precondition under which Theorem 3.12 predicts a DoRA advantage in dimension-calibrated discriminability. The realized risk bound of Theorem 3.1—which requires the magnitude-component KL omitted here—and domain-shift settings are left to future work; our experiments verify the geometric inputs to the bound, not its predictive tightness against measured test error.

A vMF Mean

The vMF partition function is $Z_d(\kappa) = \int_{\mathbb{S}^{d-1}} e^{\kappa \mu^\top x} \sigma(dx) = (2\pi)^{d/2} I_{d/2-1}(\kappa) / \kappa^{d/2-1}$. Differentiating, $\partial_\kappa \ln Z_d(\kappa) = \mathbb{E}_{\text{vMF}(\mu, \kappa)}[\mu^\top x] = A_d(\kappa)$, so $\mathbb{E}[x] = A_d(\kappa)\mu$, as used in Lemma 3.2.

Contributions

S. Zhang developed the PAC-Bayesian theory (vMF KL bounds, comparison theorem); V. Prakash implemented the fine-tuning/measurement pipeline and the four-task sweep; S. Vasanawala ran the QQP/MRPC experiments and the empirical verification. *[Placeholder—confirm/edit.]*

References

- [1] S.-Y. Liu et al., “DoRA: Weight-Decomposed Low-Rank Adaptation,” arXiv:2402.09353, 2024.
- [2] C. Si et al., “MAP: Revisiting Weight Decomposition for Low-Rank Adaptation,” arXiv:2505.23094, 2025.
- [3] E. J. Hu et al., “LoRA: Low-Rank Adaptation of Large Language Models,” arXiv:2106.09685, 2021.
- [4] D. A. McAllester, “Some PAC-Bayesian Theorems,” *Mach. Learn.*, 37(3):355–363, 1999.
- [5] Y. Liu et al., “RoBERTa: A Robustly Optimized BERT Pretraining Approach,” arXiv:1907.11692, 2019.