



Directional Updates: LoRA vs DoRA vs MAP on RoBERTa-base / MRPC



Fine-tuning metrics and aggregate directional complexity

Method	Eval acc	Trainable params	C DoRA (Σ local)	C MAP (global)	mean θ_j (rad)	mean Θ_g (rad)	mean s/d
LoRA	68.38%	1,919,234	0.65	0.005	0.0010	0.0037	0.000
DoRA	68.38%	2,002,178	0.66	0.005	0.0010	0.0038	0.000
MAP	68.38%	1,919,378	0.35	0.001	0.0008	0.0012	0.000



Θ_{global} is the single rotation of the whole flattened matrix (MAP view); θ_j are the per-column rotations (DoRA view).
Exact identity (Thm 3.10): $1 - \cos \Theta_{\text{global}} = \text{mean}_j (1 - \cos \theta_j) \rightarrow \text{DoRA sums what MAP averages.}$



Distribution of per-column directional change θ_j (pooled over all target layers)



A long right tail = a few columns rotate a lot (sparse/localized update). A concentrated body near 0 = most columns barely move.

Sparsity of directional updates and the DoRA $\leftrightarrow$ MAP identity

Per-layer sparsity ($\tau=0.01$ rad)



Where the rotation happens: attention vs MLP, and across depth

